



# Computers and systems engineering

2025

## Artificial intelligence course

Dr. Sara Khalil

### Maze-Solver-using-Reinforcement-Learning-Q-Learning

رقم الجلوس	الاسم
6358	محمد محمود محمد غانم
6349	كريم وائل طلبه محمد
6365	مصطفى حسن جاب الله عطية
6315	أنطون أنسي عبد الملك

## **Table of Contents**

### **1. Introduction**

- Problem Defined
- Objective

### **2. AI Algorithms and Tools Used**

#### **2.1. Q-learning Algorithm**

- Key Concepts
- Reward System
- Q-learning Formula
- Parameters ( $\alpha$ ,  $\gamma$ ,  $\epsilon$ )

#### **2.2. Tools and Libraries**

- Python
- NumPy
- Pygame
- Custom Gym-like Environment

### **3. Project Structure**

- main.py
- q\_learning.py
- environment.py
- visualization.py
- constants.py

### **4. Visualization & Output**

- Visualization in start
- Improvement Visualization

### **5. Conclusions**

### **6. Future Enhancements**

# 1. Introduction

This project demonstrates the application of Q-learning, a popular reinforcement learning algorithm, to solve a maze. The agent (an artificial entity) is trained to find the shortest path from a starting point to a goal within a maze. The agent learns through trial and error, receiving rewards or penalties based on its actions, and eventually converges to an optimal path.

## Problem Defined

Pathfinding in unknown environments is a classic problem in AI. Many traditional algorithms like **DFS**, BFS, or A\* require full knowledge of the environment. In contrast, **Q-learning** allows an agent to learn optimal strategies through trial-and-error interaction. The challenge lies in training the agent to avoid walls and reach the goal with the least number of moves, using only rewards and penalties.

## Objective

**Implement a Custom Maze Environment** Design a grid-based maze where the agent can interact, move, and learn.

**Apply Q-Learning Algorithm** Use the Q-learning reinforcement learning algorithm to train the agent to find the shortest path from the start point to the goal.

**Develop an Interactive Visualization** Create a real-time visual representation of the agent's movement using Pygmy to illustrate the learning result.

**Demonstrate Reinforcement Learning Concepts** Show how an agent can learn from rewards and penalties, balance exploration and exploitation, and improve its policy over time.

## 2. AI Algorithms and Tools Used

### 2.1 Q-learning Algorithm

Q-learning is a model-free reinforcement learning algorithm used to find the optimal action-selection policy for a given environment. It works by learning a **Q-table** that estimates the expected utility (**Q-values**) of taking an action in a particular state, followed by the best future actions.

#### Key Concepts:

- **State:** The position of the agent in the maze grid.
- **Action:** The movement direction (up, down, left, right).
- **Reward:** The feedback received from the environment:
  - Reaching the goal: High positive reward (e.g., +100)
  - Hitting a wall: Negative reward (e.g., -10)
  - Moving in an empty cell: Small penalty (e.g., -1)

#### Formula:

$Q(\text{state}, \text{action}) = Q(\text{state}, \text{action}) + \alpha (\text{reward} + \gamma \max(\text{Q}(\text{next\_state}, \text{all actions})) - Q(\text{state}, \text{action}))$

$$Q(s, a) \leftarrow (1 - \alpha)Q(s, a) + (\alpha) \left[ r + \gamma \max_{a'} Q(s', a') \right]$$

#### Where:

- $\alpha$ : Learning rate
- $\gamma$ : Discount factor
- $\epsilon$ : Epsilon (used in  **$\epsilon$ -greedy policy** for exploration)

## 2.2 Tools and Libraries

- **Python:** Programming language used to implement the agent and environment.
- **OpenAI Gym-like Custom Environment:** A custom environment resembling OpenAI Gym structure.
- **Pygame:** Used to create a GUI for visualizing the agent's path in the maze.
- **NumPy:** For numerical operations and Q-table storage.

## 3. Project Structure

✚ main.py	>> Main execution file
✚ q_learning.py	>> Contains the Q-learning algorithm
✚ environment.py	>> Custom environment with maze logic
✚ visualization.py	>> GUI for maze display and path visualization
✚ constants.py	>> Configuration values and parameters

## main.py

This script (main.py) is the **entry point** of a Q-learning Maze Solver project. It allows the user to choose between two maze setup modes:

1. **Randomly generated maze**
2. **Manual wall drawing using GUI**
3. **Exit to exit from program**
4. **Show or hide Training**

Then it trains an agent using Q-learning to solve the maze and finally visualizes the solution.

### 1. Import Required Modules

```
import pygame
from constants import *
from environment import MazeEnv
from q_learning import train_q_learning
from visualization import run_visualization, setup_environment
```

- pygame: GUI library used for the interface.
- constants: Assumed to store values like screen dimensions, grid sizes, colors (YELLOW, etc.).
- MazeEnv: Custom OpenAI Gym-like environment for the maze.
- train\_q\_learning: Trains the agent using the Q-learning algorithm.
- run\_visualization, setup\_environment: Show the training and environment.

### 2. Global Option Toggle

```
show_training_visualization = True
```

A global flag to show or hide training animation.

### 3. draw\_button() Function

```
def draw_button(screen, text, rect, color_normal, color_hover, mouse_pos):
```

- Draws a clickable button on the screen with hover effect.
- `rect.collidepoint(mouse_pos)` checks if the mouse is over the button.
- It then renders the button with different colors for normal/hover states.

### 4. main\_menu() Function

```
def main_menu():
```

Displays a **Pygame window** with:

- The current maze size.
- Three buttons: "Random Maze", "Manual Maze", and "Exit".
- A checkbox to toggle **training visualization**.

- **Detailed flow:**

a. Initialize Pygame

```
pygame.init()
screen = pygame.display.set_mode((800, 600))
```

b. Render UI and Buttons

```
draw_button(...) # Called for each button
```

c. Draw Checkbox

Draws a **checkbox** and a label to toggle training visualization.

If clicked, it toggles the `show_training_visualization` flag.

d. Handle Mouse Events

```
if event.type == pygame.MOUSEBUTTONDOWN:
```

If a button is clicked, return 'random', 'manual', or None.

## 5. main() Function

```
def main():
```

This function controls the program flow.

a. Show Main Menu

```
choice = main_menu()
```

Waits for user input and returns the choice.

b. Based on Choice

```
if choice == 'random':  
    env = MazeEnv(mode='random')  
    setup_environment(env)
```

- For **random**, generate and show a random maze.
- For **manual**, let the user draw walls using mouse before pressing ENTER.
- If None, user clicked "Exit".

c. Train Q-learning Agent

```
q_table = train_q_learning(env, episodes=200, epsilon_decay=0.99, min_epsilon=0.01, with_visualization=True)
```

- Trains the agent for 200 episodes with epsilon decay.
- Shows animation **if checkbox was checked**.

d. Run Final Visualization

```
run_visualization(env, q_table)
```

Displays the final agent behavior in the trained maze

## 6. Entry Point

Runs main() only if the script is executed directly

```
if __name__ == "__main__":  
    main()
```



## q\_learning.py

- Implements the core training logic.
- Initializes a Q-table of size (number of states × number of actions).
- Loops through episodes to train the agent using the epsilon-greedy strategy.
- After training, extracts the optimal path using `extract path ()` by following the maximum Q-values from the table.

### Initialize Pygame

```
pygame.init()
training_screen = pygame.display.set_mode((WINDOW_WIDTH, WINDOW_HEIGHT))
pygame.display.set_caption("Q-Learning Training Visualization")
training_clock = pygame.time.Clock()
font = pygame.font.SysFont(None, 24)
```

- Initializes the Pygame window.
- Sets up the window title, dimensions, and font for rendering text.
- A clock (`training_clock`) is used to control the frame rate during training.

### Function: `draw_training_state`

```
def draw_training_state(env, screen, font, episode, epsilon, total_reward):
```

- **Purpose:** Visually updates the screen to show the maze, agent, goal, and current stats during training.
- It does the following:
  - Clears the screen.
  - Iterates over each cell of the maze.

- If the cell is a wall (`env.maze[y][x] == 1`), it draws a **trap image**.
- Otherwise, it draws a **white cell**.
- If the cell was visited, it highlights it in **yellow**.
- Draws grid lines over the maze.
- Draws the goal as a **cheese image**.
- Draws the agent (rat) as an image.
- Refreshes the Pygame window (`pygame.display.flip()`).

### Function: `train_q_learning`

```
def train_q_learning(env, episodes=1000, alpha=0.1, gamma=0.95,
                    epsilon=0.2, epsilon_decay=0.995, min_epsilon=0.01):
```

- **Q-learning** algorithm to train the agent in the maze environment.
- Parameters:
  - alpha: Learning rate (how much to update Q-values).
  - gamma: Discount factor (importance of future rewards).
  - epsilon: Initial exploration rate (probability of taking random actions).
  - epsilon\_decay: Rate at which epsilon decreases after each episode.
  - min\_epsilon: Minimum allowed value for epsilon.

### Step-by-step:

1. **Initialize Q-table:** A table of shape `[num_states x num_actions]`, initialized to zeros.
2. Loop through each episode:
  - Reset the environment.

- Until the episode is done:
  - If a quit event is detected, quit Pygame.
  - Choose an action (**explore** or **exploit** using **epsilon-greedy policy**).
  - Execute the action and receive the next state and reward.
  - Update the Q-table using the Q-learning formula
  - Update the current state and reward.
  - Render the updated training state.
- After each episode, decay epsilon.
- Every 10 episodes, print training statistics.

### Function: `extract_path`

```
def extract_path(q_table, env, max_steps=1000):
```

- **Purpose:** Extract the path followed by the agent using the trained Q-table.
- Resets the environment and simulates steps by always choosing the best action (`argmax(q_table[state])`).
- Stops when it reaches the goal or `max_steps` is exceeded.
- Returns the path as a list of coordinates.

## environment.py

- Defines the Maze Environment class.
- Represents the maze grid, agent location, goal, available actions, and rules for state transitions.

### Class: `MazeEnv(gym.Env)`

This is a custom **reinforcement learning environment**, built by extending **`gym.Env`**.

## `__init__()` – Constructor

```
def __init__(self, mode='random'):
```

- mode='random': by default, it generates **a random maze**.
- If mode='manual', it creates an **empty grid** and lets the user draw walls.

### Key properties:

- self.observation\_space: one state per cell → GRID\_WIDTH \* GRID\_HEIGHT
- self.action\_space: 4 discrete actions: up, down, left, right
- self.agent: starts at top-left corner (0, 0)
- self.goal: bottom-right corner
- self.maze: 2D grid of 0s (free) and 1s (walls)

## Maze Initialization

### 1. Manual Maze (Empty Grid)

```
def _initialize_empty_maze(self):  
    return np.zeros((GRID_HEIGHT, GRID_WIDTH), dtype=int)
```

### 2. Random Maze

```
def _generate_maze(self):
```

- Uses a simplified **Prim's algorithm** to generate a solvable maze:
  - Start at (0,0) and mark it as free.
  - Track surrounding wall cells and randomly convert walls into paths if they only connect to one other path (to prevent cycles).

- Ensure the maze is **connected** using `_is_connected()`.

## Check for Maze Connectivity

```
def _is_connected(self, maze):
```

- Performs **Breadth-First Search (BFS)** from start to goal.
- Returns True if there's a valid path, otherwise the maze is **regenerated**.

## toggle\_wall()

```
def toggle_wall(self, x, y):
```

- Allows user to manually toggle walls in the grid (used in GUI drawing).
- Prevents toggling start and goal.

## reset()

```
def reset(self, seed=None, options=None):
```

- Resets the environment:
  - Places the agent at the start.
  - Clear the visited list.
  - Returns the initial state (cell index).

## step()

```
def step(self, action):
```

- Given an action (0=up, 1=down, 2=left, 3=right), it:
  - Moves the agent if the new position is free (0).

- Updates visited cells.
- Calculates reward:
  - Reaching goal: +10, done = True
  - Valid move: -0.1
  - Invalid (hit wall): -0.5
- Returns:
  - new\_state: cell index
  - reward
  - terminated: goal reached
  - truncated: always False here
  - Empty info {}
  -

## visualization.py

- Uses pygmy to render the maze, agent, goal, and optimal path.
- Animates the agent's movement along the learned path step-by-step.
- Highlights visited cells, the goal cell, and the path in different colors.

### 1. draw\_maze(screen, env, path=None)

is function **draws the current state of the maze** on the screen.

**Clear the screen:**

```
screen.fill(BLACK)
```

**Loop over each cell in the maze:**

```
for y in range(GRID_HEIGHT):  
    for x in range(GRID_WIDTH):
```

- maze[y][x] == 1: it's a **wall (trap)** → draw trap.png
- Else → draw **white** cell for empty space

### Highlight visited cells:

If (y, x) is in env.visited, draw a red border using:

```
pygame.draw.rect(screen, RED, rect, 1)
```

Draw grid lines (optional for clarity):

```
pygame.draw.line(...)
```

- **Draw the goal (cheese):**
  - Use cheese.png scaled to cell size, placed at the goal coordinates.
- **Draw the learned path** (if available):
  - Loop through the path and fill each cell green.
- **Draw the agent (rat):**
  - Use rat.png and draw it at env.agent.

## 2. setup\_environment(env)

This is a **GUI for manually setting walls** before training.

**Create Pygame window** with the maze size:

```
screen = pygame.display.set_mode((WINDOW_WIDTH, WINDOW_HEIGHT))
```

### 1. Handle mouse events:

- On **left-click (button 1)** → add wall (maze[y][x] = 1)

- On **right-click (button 3)** → remove wall (`maze[y][x] = 0`)

## 2. Handle keyboard events:

- If user presses Enter → exit manual setup

## 3. Redraw the maze continuously during editing.

## 3. `run_visualization(env, q_table)`

This **animates the agent solving the maze** using the Q-table.

- ❖ **Create a Pygame window** for visualization.
- ❖ **Extract the optimal path** from the Q-table:

```
path = extract_path(q_table, env)
```

Start from the initial state:

```
env.reset()
```

Step through the path one cell at a time:

- On each frame, update the agent's position.
- Draw the maze including the path so far.

Update screen at fixed speed:

```
clock.tick(10) # 10 frames per second
```

Exit if user closes the window.

Function	Purpose
<code>draw_maze</code>	Visualizes the maze grid with walls, path, and agent
<code>setup_environment</code>	Allows user to design a custom maze before training
<code>run_visualization</code>	Animates the trained agent solving the maze



## constants.py

Stores configuration values like grid size, cell size, window dimensions, and color definitions

### 1. Grid and Window Settings

```
GRID_WIDTH, GRID_HEIGHT = 6, 6
CELL_SIZE = 70
WINDOW_WIDTH = GRID_WIDTH * CELL_SIZE
WINDOW_HEIGHT = GRID_HEIGHT * CELL_SIZE
```

#### What This Does:

- Defines a **6x6 maze** grid.
- Each cell is **70 pixels × 70 pixels**.
- Total screen/window size is  $6 * 70 = 420$  pixels both in width and height.
- So the Pygame window will be **420×420 pixels**.

### 2. Visualization Frame Rate

```
Frame_RATE = 200 # FPS
```

#### What This Does:

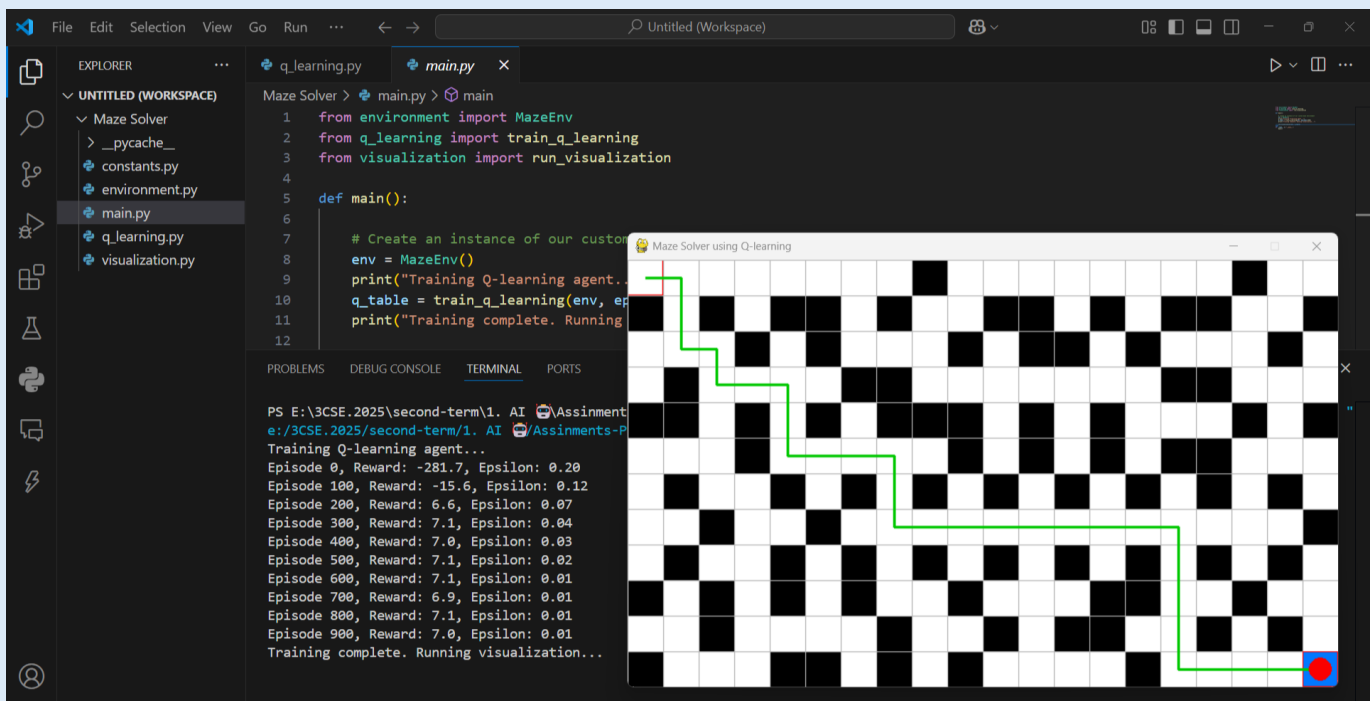
- Controls how fast the Pygame window updates.
- 200 frames per second means **very fast updates**.
- If you're doing visual training, consider lowering it (e.g., 30) to better observe the agent.

### 3. Color Definitions (RGB Format)

Variable	RGB Value	Meaning
<b>WHITE</b>	(255, 255, 255)	Empty path cells
<b>BLACK</b>	(0, 0, 0)	Background / Wall
<b>RED</b>	(250, 0, 0)	Agent (rat) border or highlight
<b>YELLOW</b>	(255, 255, 200)	Light yellow, for visited cells (unused in code above)
<b>GRAY</b>	(180, 180, 180)	Grid lines between cells
<b>BLUE</b>	(0, 120, 255)	Goal cell (optional, not used — cheese image is used instead)
<b>GREEN</b>	(0, 200, 0)	For drawing the solution path

## 4. Visualization & Output

### 4.1 Visualization in start



### Training Progress

During training, the agent gradually improved its performance. The output shows:

- **Initial episodes** had very low rewards (e.g., -281.7 at episode 0), indicating poor navigation and collisions with walls.
- As training progressed, the reward increased steadily, showing learning improvement:

Training Q-learning agent...

Episode 0, Reward: -281.7, Epsilon: 0.20

Episode 100, Reward: -15.6, Epsilon: 0.12

Episode 200, Reward: 6.6, Epsilon: 0.07

Episode 300, Reward: 7.1, Epsilon: 0.04

Episode 400, Reward: 7.0, Epsilon: 0.03

Episode 500, Reward: 7.1, Epsilon: 0.02

Episode 600, Reward: 7.1, Epsilon: 0.01

Episode 700, Reward: 6.9, Epsilon: 0.01

Episode 800, Reward: 7.1, Epsilon: 0.01

Episode 900, Reward: 7.0, Epsilon: 0.01

Training **complete**. Running visualization..

- **Epsilon decay** shows the agent relied less on exploration and more on the learned policy over time.

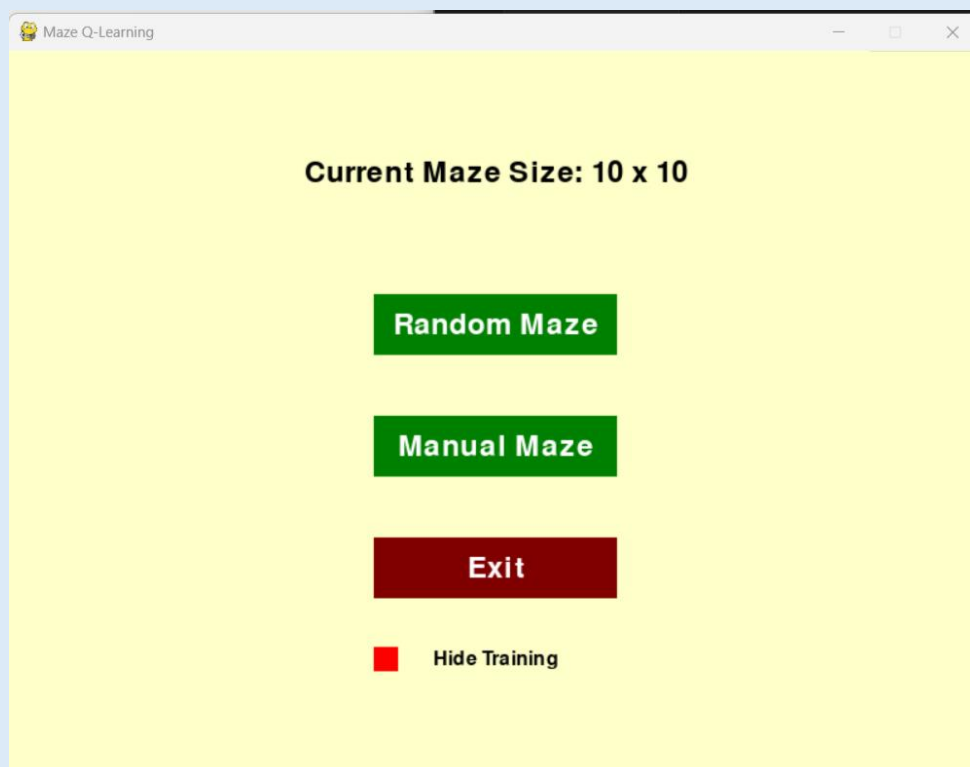
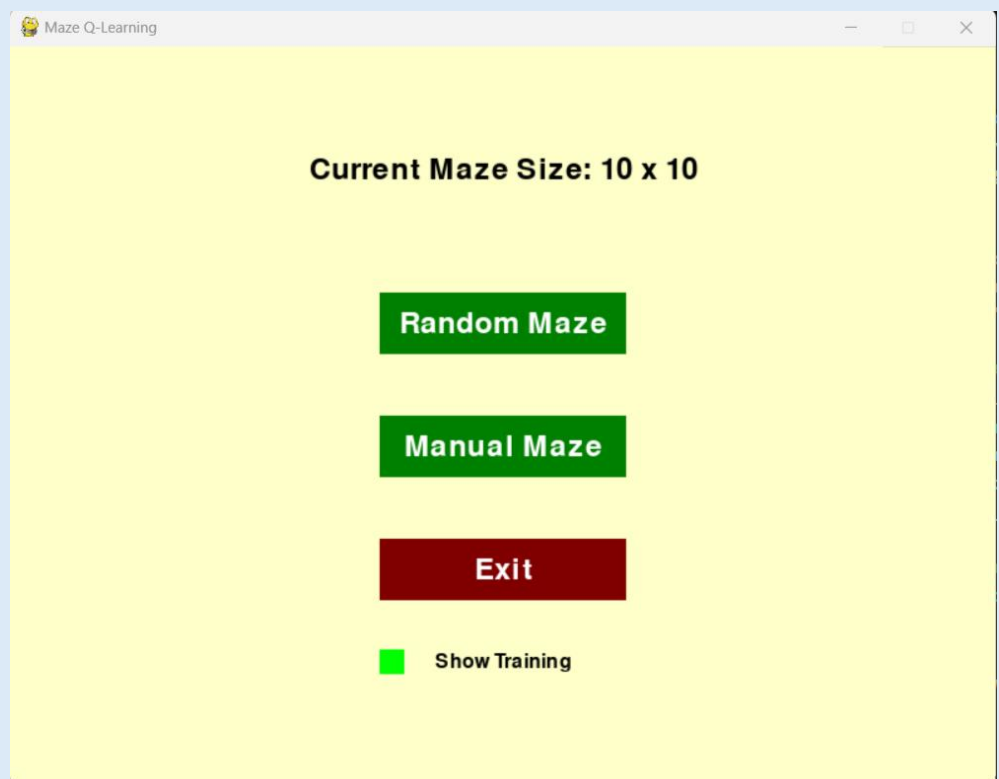
## 4.2 Final Path Visualization

The GUI shows the final path taken by the agent:

- The **green line** indicates the path taken from the **top-left (start)** to the **bottom-right (goal)**.
- **Black cells** represent walls (obstacles), and white cells are walkable.
- The **red circle** marks the goal.
- The final path is optimal and avoids all obstacles.

## 4.2 Improvement Visualization

When run code >>



### Random Maze button:

Generate randomly maze or walls

### Manual Maze button:

Generate manually maze or walls

### Exit button:

Exit from visualization

### Show Training:

To show or hide Training of agent in environment

### When selecting Show Training

#### When press on Random Maze button

Generate randomly maze or walls



**Bold Yellow** is visited cell in each Episode in training stage



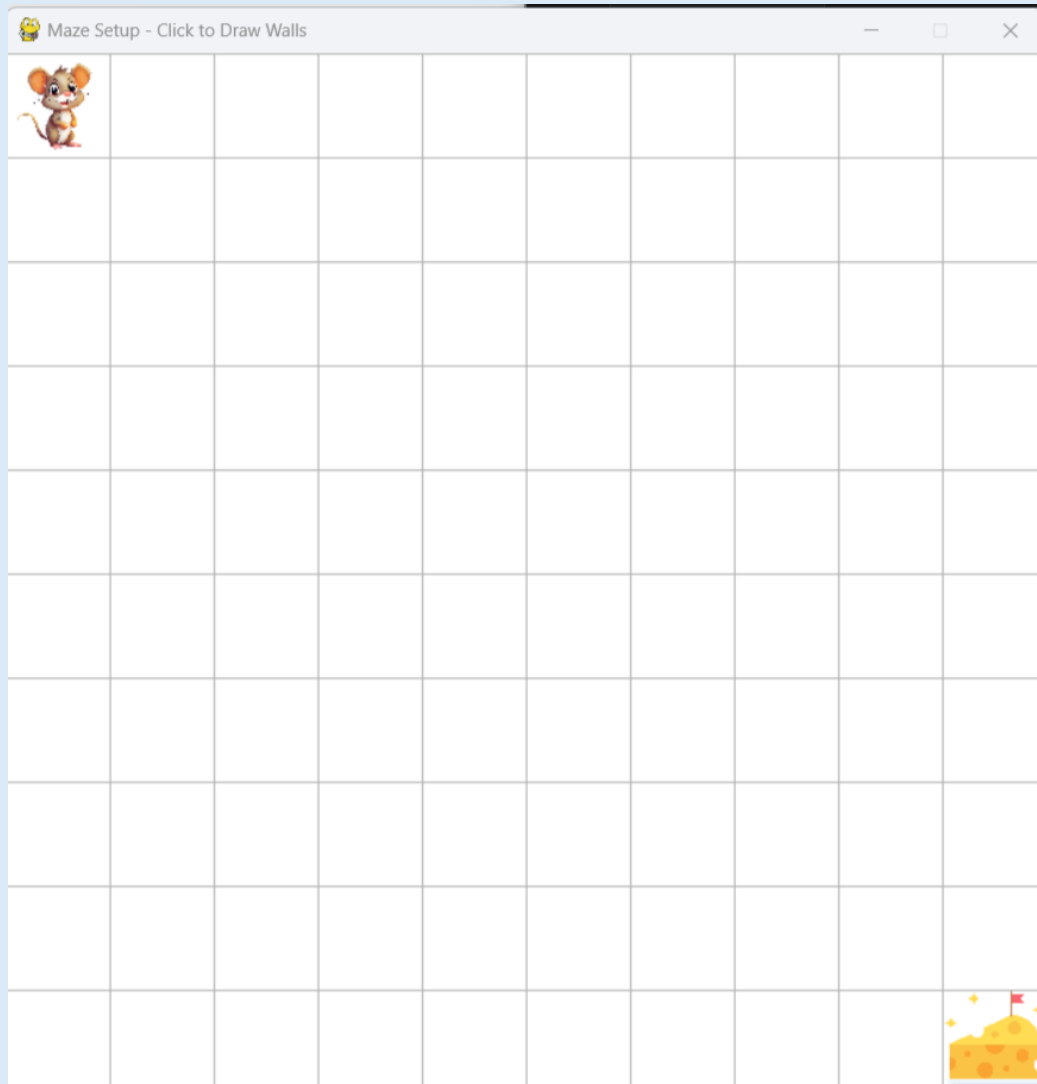
```
pygame 2.6.1 (SDL 2.28.4, Python 3.13.0)
Hello from the pygame community. https://www.pygame.org/contribute.html
User choice: random

Random maze generated automatically.
Training Q-learning agent...
Visualization enabled.
Episode: 0, Total Reward: -56.9, Epsilon: 0.20
Episode: 1, Total Reward: -177.0, Epsilon: 0.20
Episode: 2, Total Reward: -115.7, Epsilon: 0.19
Episode: 3, Total Reward: -36.9, Epsilon: 0.19
█
```

Repeat itself to get Episode : 199 , due to Episode : 200 and it started Episode : 0

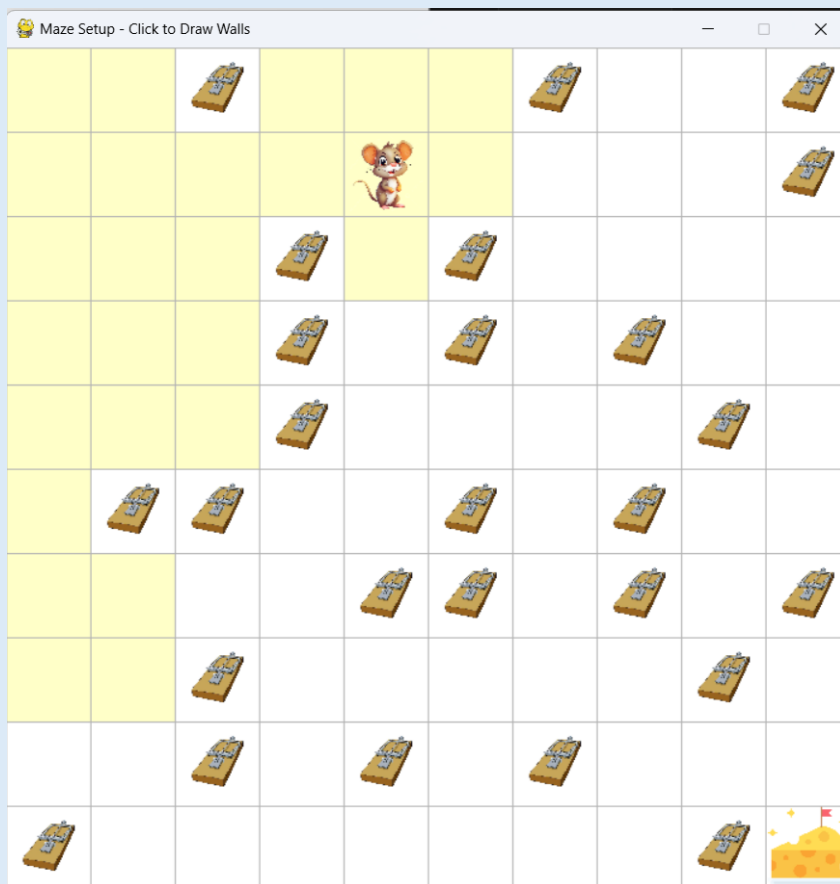
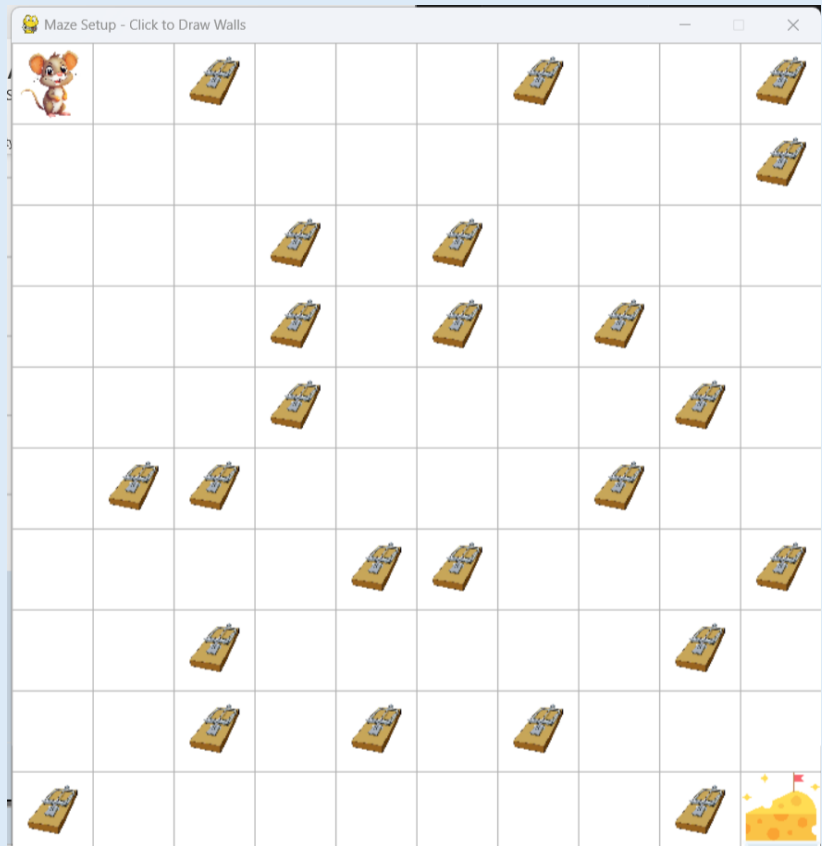
## When select Show Training

### When press on Manual Maze button



- ✓ Generate manually maze or walls
- **Right click in mouse** to add obstacle , **Left click in mouse** to remove obstacle and **Enter in keyboard** to run any training of agent!





```
Manual maze setup mode activated.  
Left-click to add walls, right-click to remove walls.  
Press ENTER to start training once done.
```

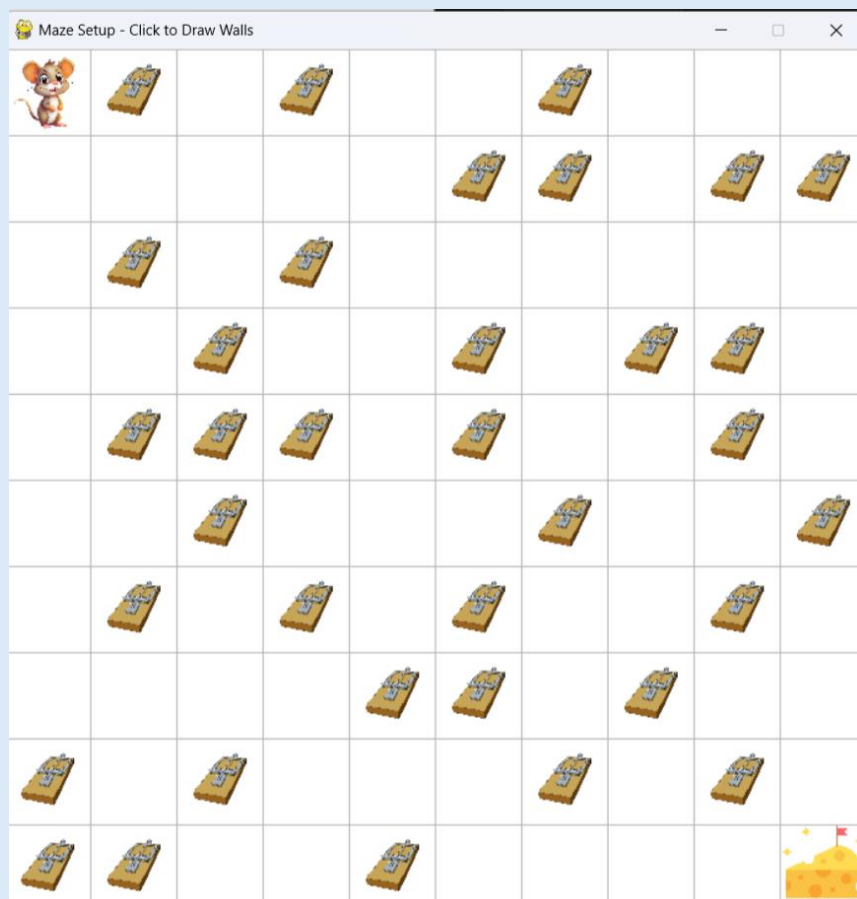
```
Training Q-learning agent...  
Visualization enabled.  
Episode: 0, Total Reward: -82.7, Epsilon: 0.20  
Episode: 1, Total Reward: -27.6, Epsilon: 0.20  
Episode: 2, Total Reward: -98.0, Epsilon: 0.19  
Episode: 3, Total Reward: -67.9, Epsilon: 0.19  
Episode: 4, Total Reward: -14.2, Epsilon: 0.19
```

Repeat to get Episode : 199 , due to Episode : 200 , it started Episode : 0

When select hide Training

When press on Random Maze button

Press on Enter to run

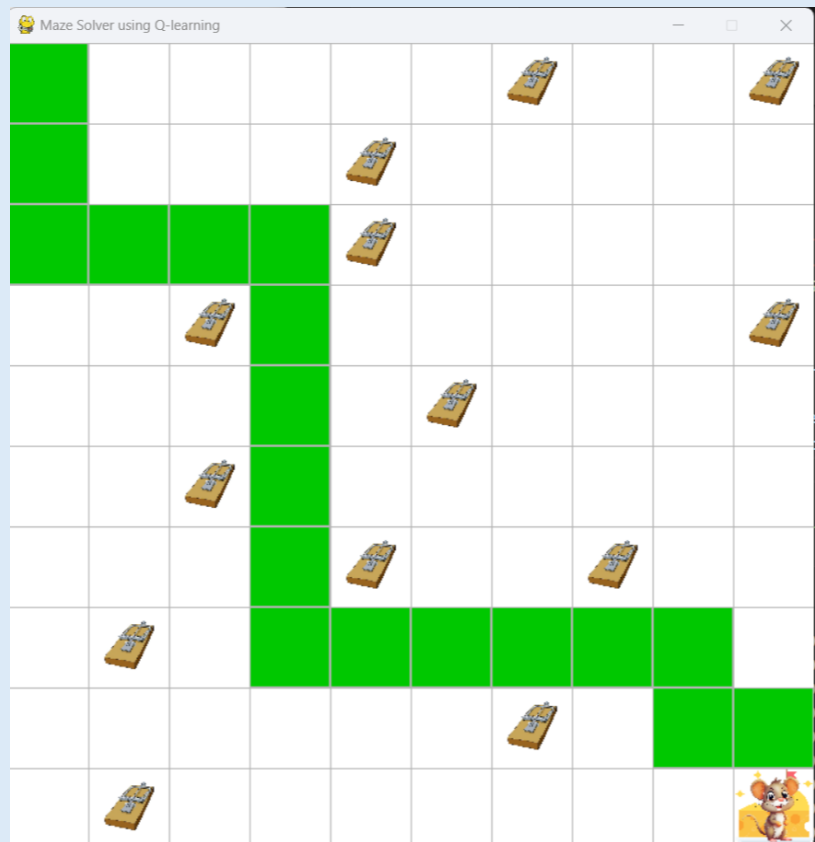




```
Episode: 191, Total Reward: 8.3, Epsilon: 0.03
Episode: 192, Total Reward: 8.3, Epsilon: 0.03
Episode: 193, Total Reward: 7.9, Epsilon: 0.03
Episode: 194, Total Reward: 8.3, Epsilon: 0.03
Episode: 195, Total Reward: 7.8, Epsilon: 0.03
Episode: 196, Total Reward: 8.2, Epsilon: 0.03
Episode: 197, Total Reward: 8.3, Epsilon: 0.03
Episode: 198, Total Reward: 8.3, Epsilon: 0.03
Episode: 199, Total Reward: 8.3, Epsilon: 0.03
Training completed.
Running visualization...
█
```

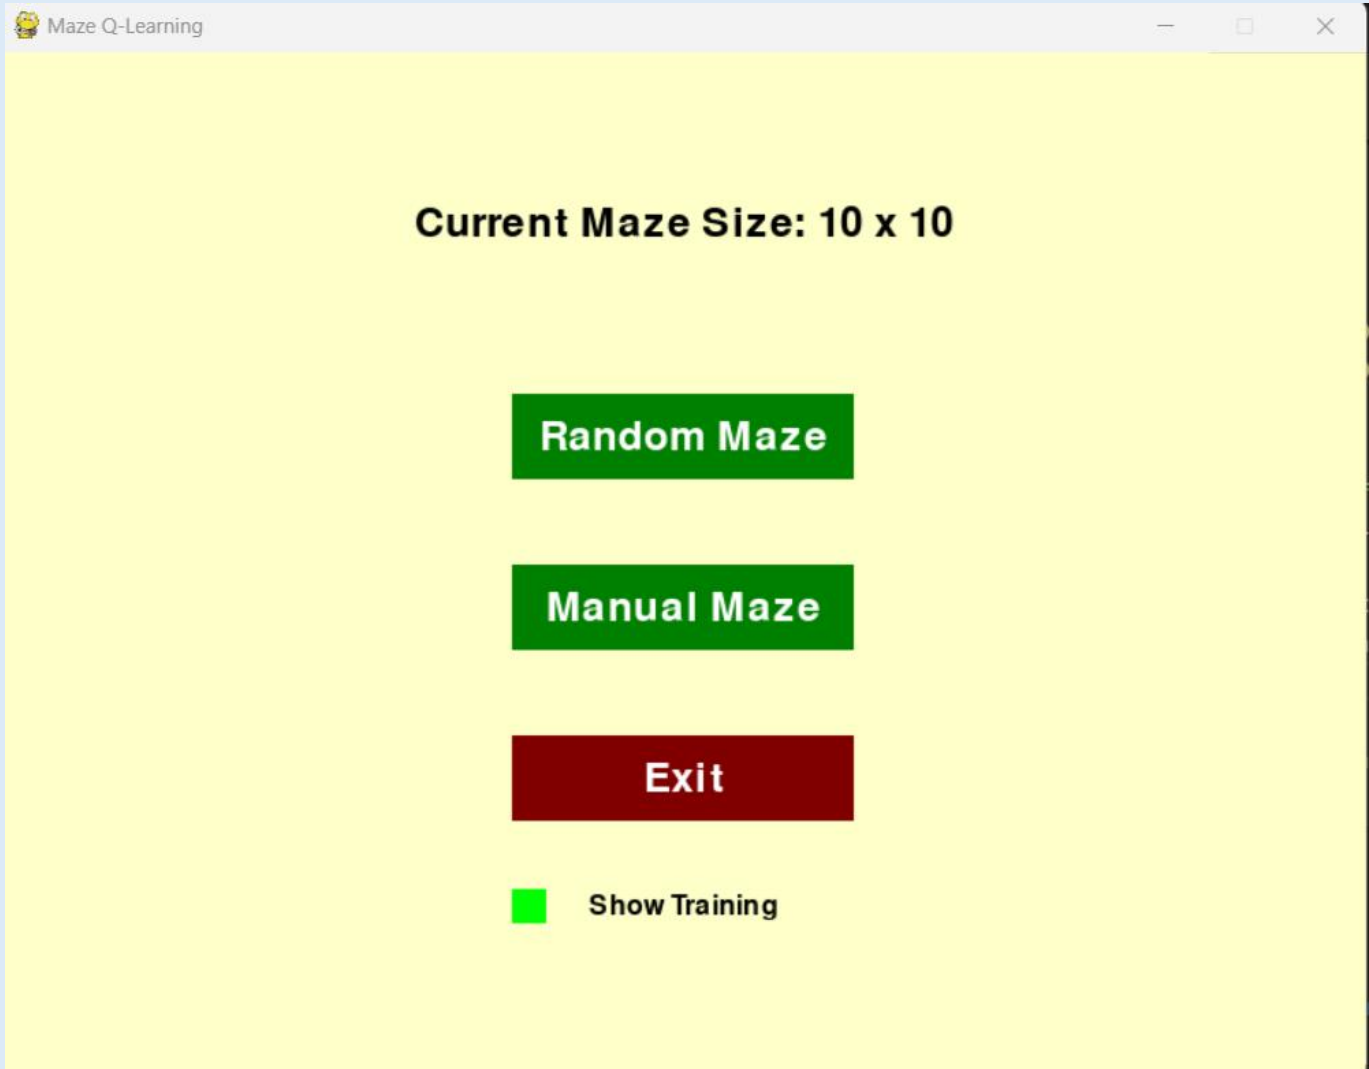
When select hide Training

When press on Manual Maze button



```
Episode: 191, Total Reward: 8.3, Epsilon: 0.03
Episode: 192, Total Reward: 8.1, Epsilon: 0.03
Episode: 193, Total Reward: 8.3, Epsilon: 0.03
Episode: 194, Total Reward: 8.3, Epsilon: 0.03
Episode: 195, Total Reward: 8.3, Epsilon: 0.03
Episode: 196, Total Reward: 7.9, Epsilon: 0.03
Episode: 197, Total Reward: 8.3, Epsilon: 0.03
Episode: 198, Total Reward: 8.3, Epsilon: 0.03
Episode: 199, Total Reward: 8.3, Epsilon: 0.03
Training completed.
Running visualization...
█
```

## To Exit



```
pygame 2.6.1 (SDL 2.28.4, Python 3.13.0)
Hello from the pygame community. https://www.pygame.org/contribute.html
User choice: None
Exiting the program.
PS E:\3CSE.2025\second-term\1. AI \Assinments-Projects\AI_Project-make-it-
realistic1\AI_Project-make-it-realistic> █
```

## 5. Conclusions

- The Q-learning algorithm effectively enables the agent to learn and navigate a complex maze.
- The training phase shows steady improvement in reward, indicating convergence.
- Visualization confirms that the learned path is valid and obstacle-free.
- The project demonstrates the capability of reinforcement learning in solving real-world pathfinding problems.

## 6. Future Enhancements

- Introduce **dynamic obstacles** for advanced learning.
- Use **deep reinforcement learning (DQN)** for larger and more complex mazes.
- Implement **multi-agent** scenarios for cooperative maze solving.